Architecture: how the agent works, start to end

Answer path (top), participants (middle), evaluation path (bottom)

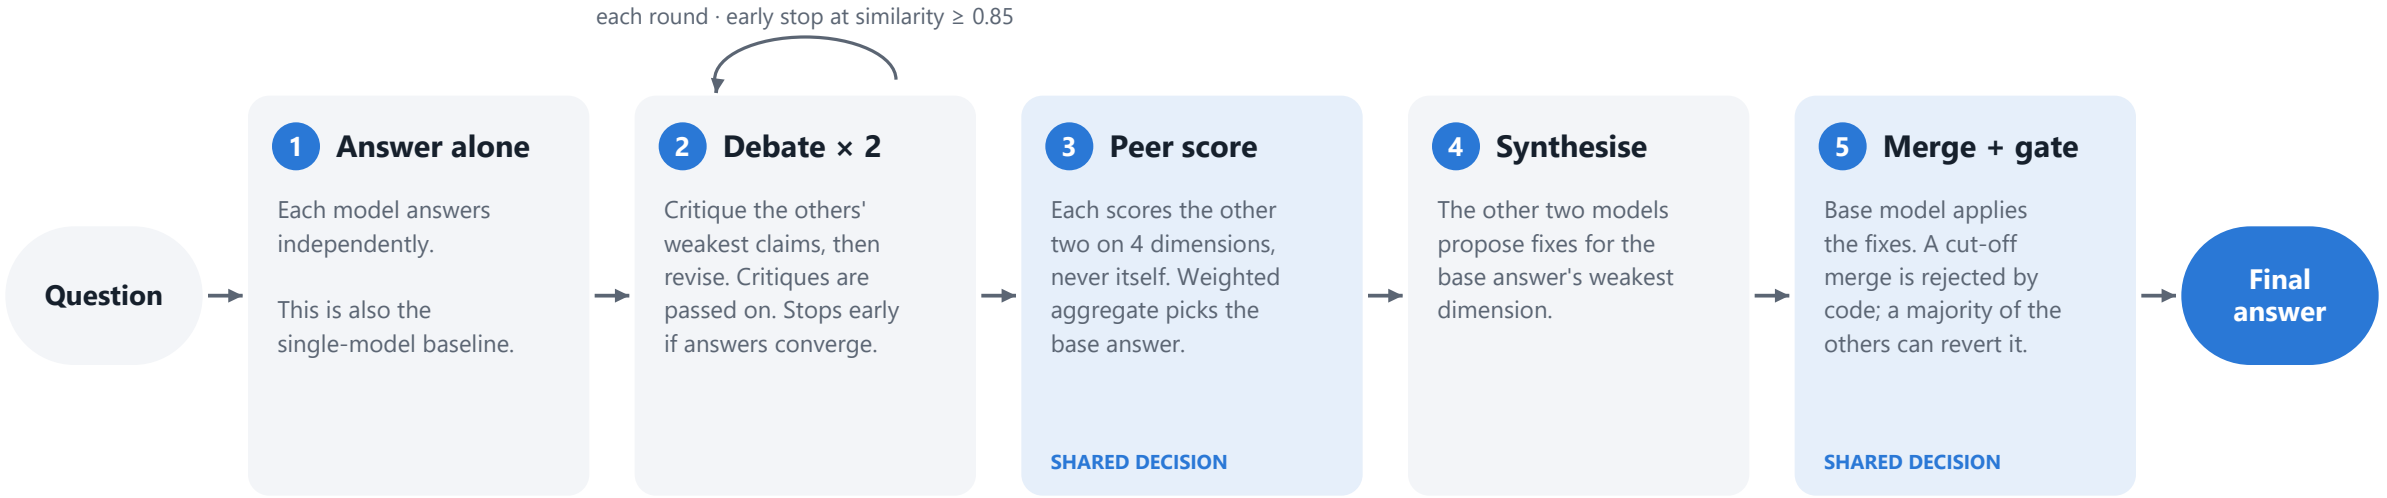

Debaters Llama 3.2 3B · Meta | Qwen 2.5 3B · Alibaba | Gemma 3 4B · Google — 3 models from 3 vendors, run via ollama

Evaluation — proving it was worth doing (runs separately; not part of producing the answer)

Baseline	Debate answer	Independent judges	Verdict
each model's step-1 answer	the final answer above	Phi-4 mini (Microsoft) + Ministral 3 3B (Mistral)	debate vs best single model

No judge shares a vendor with a debater. The judges are tested too: deliberately damaged answers must score lower (slide 2).

Guardrails: hard $25 spend ceiling in code · keys from environment variables only · one command to run · dry-run mode at zero cost · full transcript of every exchange